

Coursework 1: Image filtering

The coursework includes coding questions and/or written questions. Please read both the text and the code in this notebook to get an idea what you are expected to implement.

What is expected?

- Complete and run the code using `jupyter-lab`.
- Export (File | Save and Export Notebook As...) the notebook as a PDF file, which contains your code, results and answers, and upload the PDF file onto [Scientia](#).
- Instead of clicking the Export button, you can also run the following command instead: `jupyter nbconvert coursework_01.ipynb --to pdf`
- If Jupyter complains issues during exporting, it is likely that [pandoc](#) or [latex](#) is not installed, or their paths have not been included. You can install the relevant libraries and retry. Alternatively, use the Print function of your browser to export the PDF file.
- If Jupyter-lab does not work for you at the end, alternatively, you can use Google Colab to write the code and export the PDF file.

Dependencies:

You may need to install [Jupyter-Lab](#) and other libraries used in this coursework, such as by running the command: `pip3 install [package_name]`

```
In [1]: # Import libraries (provided)
import imageio.v3 as imageio
import numpy as np
import matplotlib.pyplot as plt
import scipy
import scipy.signal
import math
import time
```

Q1. Moving average filtering (20 points).

Read the provided clean image, add noise to the image and design a moving average filter for denoising.

You are expected to design the kernel of the filter and then perform 2D image filtering using the function `scipy.signal.convolve2d()`.


In [2]:

```
# Read the image (provided)
image = imageio.imread('campus_snow.jpg')

# Corrupt the image with Gaussian noise (provided)
noise_mu = 0
noise_sigma = 50
noise = np.random.normal(noise_mu, noise_sigma, image.shape)
image_noisy = image + noise

# Visualise the images (provided)
plt.subplot(1, 2, 1)
plt.imshow(image, cmap='gray')
plt.title('Clean image')

plt.subplot(1, 2, 2)
plt.imshow(image_noisy, cmap='gray')
plt.title('Noisy image')
plt.gcf().set_size_inches(8, 4)
```



Q1.1 Filter the noisy image using a 5x5 moving average filter. Display the filtered image.

In [3]:


```
# Design the filter h
### Insert your code ####
# Creating 5x5 matrix, with each cell = 1/25
h = np.full((5,5), 1/25)

# Convolve the noisy image with h using scipy.signal.convolve2d function
### Insert your code ####
# Symmetrically reflecting the pixel values at the boundary when filtering
image_filtered = scipy.signal.convolve2d(image_noisy, h, mode='same', boundary='symm')

# Print the filter (provided)
print('Filter h = {0}'.format(h))

# Display the filtering result (provided)
plt.imshow(image_filtered, cmap='gray')
plt.title('After moving average filtering')
plt.gcf().set_size_inches(4, 4)
```

```
Filter h = [[0.04 0.04 0.04 0.04 0.04]
[0.04 0.04 0.04 0.04 0.04]
[0.04 0.04 0.04 0.04 0.04]
[0.04 0.04 0.04 0.04 0.04]
[0.04 0.04 0.04 0.04 0.04]]
```



Q1.2 Assess the quality of the filtered image using a quantitative metric, the peak signal-to-noise ratio (PSNR).

For this case, the pixel intensity of the image is represented using the uint8 format, with the peak value to be 255. The PSNR is defined as,

$$\text{PSNR} = 10 \cdot \log_{10} \left(\frac{255^2}{\frac{1}{N} \sum_x [J(x) - I(x)]^2} \right)$$

where x denotes the pixel index, N denotes the total number of pixels in the image, J denotes the filtered i.e. denoised image and I denotes the ground truth clean image. The denominator of the term within the logarithm operator is the mean squared error between I and J .

You can find more detail about PSNR [here](#).

```
In [4]: # Implement the PSNR function
def eval_psnr(I, J):
    # I: the ground truth clean image (peak value: 255 for uint8 data format)
    # J: the denoised image
    #
    # return: the PSNR metric (unit: dB)

    ### Insert your code ####
    mse = np.sum((I - J)**2) / I.size
    psnr = 10 * math.log10(255**2 / mse)
    return psnr

# Evaluate the PSNR for the filtered image (provided)
psnr = eval_psnr(image, image_filtered)

# Print the PSNR (provided)
print('PSNR = {0:.2f} dB'.format(psnr))
```

PSNR = 18.49 dB

Q2. Gaussian filtering (70 points).

Q2.1 Implement a function that constructs a 2D Gaussian filter given the parameter σ .

```
In [5]: # Implement the Gaussian filter
def gaussian_filter_2d(sigma):
    # sigma: the parameter sigma for the Gaussian kernel (unit: pixel)
    #
    # return: a 2D array for the Gaussian kernel

    # The filter radius is 4 times sigma (provided)
    rad = int(math.ceil(4 * sigma))
    sz = 2 * rad + 1

    # Calculate the filter weights
    ### Insert your code ####
    # Creating 2D array centred on (0,0)
    range_1d = np.arange(-rad, rad + 1)
    x, y = np.meshgrid(range_1d, range_1d)
```


```

# Applying Gaussian formula
h = np.exp(-(x**2 + y**2) / (2 * sigma**2))

# Normalise by the sum (rather than the analytic  $2 * \pi * \sigma^2$ )
h /= h.sum()
return h

# Visualise the Gaussian filter when  $\sigma = 5$  pixel (provided)
sigma = 5
h = gaussian_filter_2d(sigma)
plt.imshow(h)
plt.title('Gaussian kernel')
plt.gcf().set_size_inches(4, 4)

```



Q2.2 Perform Gaussian filtering ($\sigma = 5$ pixels) for the noisy image, evaluate the computational time for Gaussian filtering and display the filtered image.

```
In [6]: # Construct the Gaussian filter (provided)
```


```
sigma = 5
h = gaussian_filter_2d(sigma)

# Perform Gaussian filtering and count time
### Insert your code ####
start = time.time()
# Using symmetric pixel values for filtering at the edge of the image
image_smoothed_2d = scipy.signal.convolve2d(image_noisy, h, mode='same', boundary='symm')

elapsed = time.time() - start
print("Time taken: {0:.4f} seconds".format(elapsed))

# Visualise the filtered image (provided)
plt.imshow(image_smoothed_2d, cmap='gray')
plt.title('Gaussian smoothed image')
plt.gcf().set_size_inches(4, 4)
```

Time taken: 1.5237 seconds



Q2.3 Implement a function that generates a 1D Gaussian filter given the parameter σ . Construct 1D Gaussian filters along x-axis and y-axis respectively.

In [7]:

```
# Implement the 1D Gaussian filter
def gaussian_filter_1d(sigma):
    # sigma: the parameter sigma in the Gaussian kernel (unit: pixel)
    #
    # return: a 1D array for the Gaussian kernel

    # The filter radius is 4 times sigma (provided)
    rad = int(math.ceil(4 * sigma))
    sz = 2 * rad + 1

    # Calculate the filter weights
    ### Insert your code ####
    # First, generating an array of numbers of length sz centred on 0
    array = np.arange(-rad, rad+1)

    # Create 1D array using Gaussian formula
    h = np.exp(-array**2 / (2 * sigma**2))

    # Normalise by the sum (rather than the analytic  $\sqrt{2\pi} \sigma$ )
    h /= h.sum()
    return h


# sigma = 5 pixel (provided)
sigma = 5

# Construct the Gaussian filter along x-axis. Its shape is (1, sz).
### Insert your code ####
h = gaussian_filter_1d(sigma)
h_x = h.reshape(1, -1) # -1 inferred to be sz

# Construct the Gaussian filter along y-axis. Its shape is (sz, 1).
### Insert your code ####
h_y = h.reshape(-1, 1)

# Visualise the filters (provided)
plt.subplot(1, 2, 1)
plt.imshow(h_x)
```

```
plt.subplot(1, 2, 2)
plt.imshow(h_y)
plt.gcf().set_size_inches(4, 4)
```



Q2.4 Perform Gaussian filtering ($\sigma = 5$ pixels) using two separable filters and evaluate the computational time for separable Gaussian filtering. Compare the smoothed image using separable filtering to the smoothed image using a single 2D Gaussian filter.

In [8]:

```
# Perform separable Gaussian smoothing and count time
### Insert your code ####
start = time.time()
# First, filtering along the x-axis
# You could also use scipy.ndimage.convolve1d if you used h directly as a 1D array
# Rather than the reshaped h_x and h_y arrays.
image_smoothed = scipy.signal.convolve2d(image_noisy, h_x, mode='same', boundary='symm')
# Now along y-axis
image_smoothed = scipy.signal.convolve2d(image_smoothed, h_y, mode='same', boundary='symm')

elapsed = time.time() - start
print("Time taken: {:.4f} seconds".format(elapsed))
```

```
# Report the difference between the separably filtered image and the image filtered by a single 2D Gaussian filter (provided)
diff = image_smoothed - image_smoothed_2d
print('Mean absolute difference = {:.6f}'.format(np.mean(np.abs(diff))))
```

Time taken: 0.2007 seconds
Mean absolute difference = 0.000000

Q2.5 Perform Gaussian smoothing for the same noisy image, assess the quality of the Gaussian smoothed image using PSNR, when different sigma values are used.

```
In [9]: # A list of sigma values (provided)
list_sigma = np.arange(0.5, 5, 0.5)


# Perform Gaussian smoothing with different sigma values and record the PSNR values
### Insert your code ####
list_psnr = []
for sigma in list_sigma:
    h = gaussian_filter_1d(sigma)
    h_x = h.reshape(1, -1)
    h_y = h.reshape(-1, 1)

    # Separable filtering
    image_smoothed = scipy.signal.convolve2d(image_noisy, h_x, mode='same', boundary='symm')
    image_smoothed = scipy.signal.convolve2d(image_smoothed, h_y, mode='same', boundary='symm')

    list_psnr.append(eval_psnr(image, image_smoothed))

# Plot the PSNR metric against sigma (provided)
plt.plot(list_sigma, list_psnr)
plt.xlim([0, 5])
plt.xlabel('$\sigma$ (unit: pixel)')
plt.ylabel('PSNR (unit: dB)')
plt.grid()
plt.gcf().set_size_inches(4, 4)
```

```
<>:22: SyntaxWarning: "\s" is an invalid escape sequence. Such sequences will not work in the future. Did you mean "\\\s"? A raw string is also an option.  
<>:22: SyntaxWarning: "\s" is an invalid escape sequence. Such sequences will not work in the future. Did you mean "\\\s"? A raw string is also an option.  
/var/folders/pt/5mg57l7j27b2bbpf3wt5bm540000gn/T/ipykernel_79377/1282904532.py:22: SyntaxWarning: "\s" is an invalid escape sequence. Such sequences will not work in the future. Did you mean "\\\s"? A raw string is also an option.  
    plt.xlabel('$\sigma$ (unit: pixel)')
```



Q2.6 Implement 3x3 Sobel filters, perform Sobel filtering for the noisy image, and display the gradient magnitude map.

In [10]:

```
# Construct the Sobel filters  
### Insert your code ###  
sobel_x = np.array([[1,0,-1],  
                   [2,0,-2],  
                   [1,0,-1]])  
sobel_y = sobel_x.T  
  
# Print the filters (provided)
```

```

print('Sobel_x = {0}'.format(sobel_x))
print('Sobel_y = {0}'.format(sobel_y))

# Sobel filtering for the noisy image
def sobel_filter(image: np.ndarray):
    # Note: I separate the x- and y-axis Sobel filters into 1D vectors to speed up the convolutions
    # Initialising the separable row vectors (with shape (1,3))
    weight_vector = np.array([[1,2,1]])
    diff_vector = np.array([[1,0,-1]])

    # First I apply the Sobel filter along the x axis (i.e. getting the gradients with respect to the x-axis)
    g_x = scipy.signal.convolve2d(image, weight_vector.T, mode='same', boundary='symm')
    g_x = scipy.signal.convolve2d(g_x, diff_vector, mode='same', boundary='symm')

    # Now filtering along the y axis (note the vectors are different)
    g_y = scipy.signal.convolve2d(image, diff_vector.T, mode='same', boundary='symm')
    g_y = scipy.signal.convolve2d(g_y, weight_vector, mode='same', boundary='symm')

    return g_x, g_y

g_x, g_y = sobel_filter(image_noisy)

# Calculate the gradient magnitude
### Insert your code ####
grad_mag_noisy = np.sqrt(g_x**2 + g_y**2)


# Display the magnitude map (provided)
plt.imshow(grad_mag_noisy, cmap='gray', vmin=0, vmax=500)
plt.title('Gradient magnitude map')
plt.gcf().set_size_inches(4, 4)

```

```

Sobel_x = [[ 1  0 -1]
           [ 2  0 -2]
           [ 1  0 -1]]
Sobel_y = [[ 1  2  1]
           [ 0  0  0]
           [-1 -2 -1]]

```



Q2.7 Perform Gaussian smoothing for the noisy image, followed by Sobel filtering and display the gradient magnitude map.

In [11]:

```
# Parameter for the Gaussian filter (provided)
sigma = 5

# Gaussian smoothing
### Insert your code ###

# Copied this code from Q2.4 (using 2 separable filters)
# Smoothing along x-axis
image_smoothed = scipy.signal.convolve2d(image_noisy, h_x, mode='same', boundary='symm')
# Now along y-axis
image_smoothed = scipy.signal.convolve2d(image_smoothed, h_y, mode='same', boundary='symm')

# Sobel filtering
### Insert your code ###
```

```


smooth_g_x, smooth_g_y = sobel_filter(image_smoothed)

# Calculate the gradient magnitude
### Insert your code ####
grad_mag = np.sqrt(smooth_g_x**2 + smooth_g_y**2)

# Display the magnitude map (provided)
plt.imshow(grad_mag, cmap='gray', vmin=0, vmax=100)
plt.title('Gradient magnitude map after Gaussian smoothing')
plt.gcf().set_size_inches(4, 4)

```

Gradient magnitude map after Gaussian smoothing



Q3. Implement image filters using Pytorch (10 points).

[Pytorch](#) is a machine learning framework that supports filtering and convolution.

The [Conv2D](#) operator takes an input array of dimension $N \times C_1 \times X \times Y$, applies the filter and outputs an array of dimension $N \times C_2 \times X \times Y$. Here, since we only have one image with one colour channel, we will set $N=1$, $C_1=1$ and $C_2=1$. You can read the documentation of Conv2D for more detail.

```
In [12]: # Import libraries (provided)
import torch
```

Q3.1 Expand the dimension of the noisy image into 1x1xXxY and convert it to a Pytorch tensor.

```
In [13]: # Expand the dimension of the numpy array
### Insert your code ####
expanded_image = np.expand_dims(image_noisy, axis=(0, 1))

# Convert to a Pytorch tensor using torch.from_numpy
### Insert your code ####

# I convert it to float32 to speed up computational time (it took ages on my computer for some reason)
tensor_noisy = torch.from_numpy(expanded_image).float()
```

Q3.2 Create a Pytorch Conv2D filter, set its kernel to be a 2D Gaussian filter, perform filtering, report computational time and display the result.

```
In [14]: # A 2D Gaussian filter when sigma = 5 pixel (provided)
sigma = 5
h = gaussian_filter_2d(sigma)

# Construct the Conv2D filter
### Insert your code ####
# Converting to float32 as in Q3.1
filter = torch.from_numpy(np.expand_dims(h, axis=(0, 1))).float()

# Filtering and assess computational time
### Insert your code ####
start = time.time()

# Zero-padding so the resulting image is the same size as the original
image_filtered = torch.nn.functional.conv2d(tensor_noisy, filter, padding='same')

elapsed = time.time() - start
print("Time taken: {0:.4f} seconds".format(elapsed))

# squeezing tensor back to 2D
image_filtered = image_filtered.squeeze().numpy()
```

```
# Display the filtering result (provided)
plt.imshow(image_filtered, cmap='gray')
plt.title('Gaussian smoothed image using PyTorch')
plt.gcf().set_size_inches(4, 4)
```

Time taken: 37.4893 seconds

Gaussian smoothed image using PyTorch

